

Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting

arXiv: 1912.09363v1

[Read the full paper here](#)

AI-generated keywords: Temporal Fusion Transformer

AI-generated Summary

The paper "Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting" by Bryan Lim, Sercan O. Arik, Nicolas Loeff, and Tomas Pfister addresses the challenge of multi-horizon forecasting problems that involve a complex mix of inputs, including static covariates, known future inputs, and exogenous time series observed historically without prior information on how they interact with the target. While deep learning models have been proposed for multi-step prediction, they often lack interpretability and do not account for the full range of inputs present in common scenarios. To address these issues, the authors introduce the Temporal Fusion Transformer (TFT), an attention-based architecture that combines high-performance multi-horizon forecasting with interpretable insights into temporal dynamics. The TFT utilizes recurrent layers for local processing and interpretable self-attention layers to learn long-term dependencies at different scales. The model also includes specialized components for selecting relevant features and gating layers to suppress unnecessary components. This enables high performance in a wide range of regimes while providing practical interpretability use cases. The authors demonstrate significant performance improvements over existing benchmarks on various real-world datasets using TFT. They showcase three practical interpretability use cases of TFT: feature importance analysis, anomaly detection, and counterfactual analysis. Overall, this paper presents a novel approach to multi-horizon time series forecasting that combines high accuracy with interpretability and has potential applications in various domains such as finance, healthcare, and transportation.

AI-generated Key Points

- The paper addresses the challenge of multi-horizon forecasting problems that involve a complex mix of inputs.
- Deep learning models for multi-step prediction often lack interpretability and do not account for the full range of inputs present in common scenarios.
- The authors introduce the Temporal Fusion Transformer (TFT), an attention-based architecture that combines high-performance multi-horizon forecasting with interpretable insights into temporal dynamics.
- The TFT utilizes recurrent layers for local processing and interpretable self-attention layers to learn long-term dependencies at different scales.
- The model also includes specialized components for selecting relevant features and gating layers to suppress unnecessary components, enabling high performance in a wide range of regimes while providing practical interpretability use cases.
- The authors demonstrate significant performance improvements over existing benchmarks on various real-world datasets using TFT.
- Three practical interpretability use cases of TFT are showcased: feature importance analysis, anomaly detection, and counterfactual analysis.
- Overall, this paper presents a novel approach to multi-horizon time series forecasting that combines high accuracy with interpretability and has potential applications in various domains such as finance, healthcare, and transportation.

AI-generated Lay Summary

This paper talks about a problem with predicting things that happen in the future. Sometimes, it's hard to understand why the prediction is made. The authors made a new way to predict things called Temporal Fusion Transformer (TFT). It can predict things accurately and also explain why it makes those predictions. They tested it on real-world data and it worked better than other ways of predicting. There are three ways to use TFT: finding important features, detecting unusual events, and figuring out what could have happened differently.

- Multi-horizon forecasting: predicting things that will happen in the future over different time periods.
- Deep learning models: computer programs that learn from data to make predictions or decisions.
- Interpretable: easy to understand or explain.
- Temporal dynamics: how things change over time.
- Recurrent layers: parts of a computer program that remember past information.
- Self-attention layers: parts of a computer program that focus on important information within itself.
- Features: characteristics or attributes used for prediction.
- Gating layers: parts of a computer program that control the flow of information by deciding which parts are important and which aren't.
- Benchmarks: standards used for comparison with other methods or systems.

AI-generated Blog-like Article

Temporal Fusion Transformers for Interpretable Multi-Horizon Time Series Forecasting

Time series forecasting is a challenging problem that requires the ability to accurately predict future values based on past observations. Traditional methods such as linear regression and autoregressive models are limited in their ability to capture complex temporal dynamics, making them inadequate for many real-world applications. Deep learning models have been proposed as an alternative, but they often lack interpretability and do not account for the full range of inputs present in common scenarios. In this paper, Bryan Lim, Sercan O. Arik, Nicolas Loeff, and Tomas Pfister propose the Temporal Fusion Transformer (TFT), an attention-based architecture that combines high performance multi-horizon forecasting with interpretable insights into temporal dynamics.

Background

Time series forecasting involves predicting future values based on past observations of a given variable or set of variables over time. This type of prediction can be used in various domains such as finance, healthcare, and transportation to make decisions about investments or operations management. However, traditional methods such as linear regression and autoregressive models are limited in their ability to capture complex temporal dynamics due to their reliance on fixed weights and assumptions about stationarity. As a result, these methods often fail when applied to real-world problems with nonlinear relationships between input variables or multiple sources of information at different scales. Deep learning models have emerged as an alternative approach for time series forecasting due to their superior performance compared to traditional methods. These models use neural networks with multiple layers of neurons connected by weights that can be adjusted through training data sets using backpropagation algorithms. While deep learning has proven effective in many cases, it lacks interpretability which makes it difficult to understand how the model arrives at its predictions or identify potential errors in its output without extensive manual analysis or trial-and-error experimentation with hyperparameters settings. Furthermore, existing deep learning architectures do not account for all types of inputs present in common scenarios such as static covariates (elements that remain constant over time) known future inputs (events whose occurrence is known ahead of time), and exogenous time series observed historically without prior information on how they interact with the target variable being predicted).

The Temporal Fusion Transformer Model

To address these issues related to deep learning approaches for multi-step prediction tasks involving multiple sources of input data at different scales, the authors introduce TFT – a novel attention based architecture combining high performance multi horizon forecasting capabilities along with interpretability into temporal dynamics. The TFT utilizes recurrent layers for local processing while self-attention layers are employed learn long term dependencies across different scales. Additionally, specialized components like feature selection modules & gating layers help suppress unnecessary components from being considered during inference. This enables better accuracy & improved performance across various regimes while providing practical interpretability use cases.

Experimental Results

The authors demonstrate significant improvements over existing benchmarks on various real world datasets using TFT. They showcase three practical interpretability use cases: feature importance analysis, anomaly detection & counterfactual analysis. Feature importance analysis helps identify important features contributing towards predictions made by TFT while anomaly detection helps detect any unexpected behavior within the dataset which could lead further investigation into possible causes behind it. Counterfactual analysis allows users examine what would happen if certain conditions were changed thus helping gain deeper insights into underlying system behavior under varying circumstances.

Conclusion

Overall , this paper presents a novel approach towards multi horizon timeseries forecasting combining both accuracy & interpretability together having potential applications across various domains like finance , healthcare & transportation etc.. The authors successfully demonstrate significant improvement over existing benchmarks using TFT along with showcasing 3 practical use cases where it can be applied effectively giving rise new possibilities within field machine learning research especially related timeseries prediction tasks

From SummarizePaper.com (2026. 3. 18. 12:30:55)